

LATENT OCEAN

Data that organizes itself.

Diren Kumaratilleke · direnavk@outlook.com

www.latent-ocean.com

THE PROBLEM

Today's data systems chase what's common and miss what matters. The most valuable patterns hide in the rare events, the unusual signals, the connections nobody is looking for.

WHY NOW

Data is surfacing like never before; it deserves vector alignment for maximum signal at scale.

The next generation of data interpretation has to ingest its own structure.

INGESTION

Read any kind of data which can be structured or messy, big or small.

FOLLOW THE FLOW

Don't just see what's there and be capable of tracking how data patterns move.

REORGANIZE ITSELF

Build its own structure as it learns. No engineer required.

WHAT IT IS

A new way to read data;
one that lets every kind of information speak the
same language.

Lets every kind of data line up against every other.

LESS NOISE

REUSABLE PARTS

ONE SHARED
LANGUAGE

WHAT MAKES IT DIFFERENT

Most data systems are designed by hand. Latent Ocean grows on its own with no setup needed.

TODAY'S DATA

- Designed once
- Architecture frozen
- Domain-specific
- Brittle to new data

LATENT OCEAN

- Discovers its own structure
- Reorganizes itself
- Works across domains
- Grows with new data

WHERE IT APPLIES

One system. Many worlds.



FINANCE

Reading the SEC archive as one connected story and surfacing the companies that don't fit.



MEDICINE

Finding the connections across millions of medical papers that no practitioner has time to read.



RESEARCH

Connecting centuries of scientific thinking — finding the threads no one person could see in a lifetime of reading.

AND TOOLS BUILT ON TOP — FOR EVERY OTHER FIELD

Real findings. Real domains.

1

Found rare, high-value patterns in messy real-world data — patterns that line up with what experts already know, and reveal what they don't.


2

Reads live data and surfaces what matters most — adjustable by region and timeframe.

3

Applied a single system to multiple entirely different fields, without retraining.

Every row, structurally aware.



The home page — what visitors see first.

Structural fingerprints. A first-class column.

THE PRIMITIVE · A DATA TYPE, NOT A PRODUCT

Structural fingerprints. A first-class column.

Every row that passes through Latent Ocean acquires two deterministic columns: a **48-bit structural fingerprint** and a **4-dimensional score vector**. The fingerprint locates the row in a compressed representation of its universe. The score vector reports how structurally unusual that location is. Both are reproducible at bit level under a fixed seed.

What it is

THE FINGERPRINT

48 bits. Deterministic rotation-ensemble hash of the row's structured attributes. Two rows whose signatures are close in Hamming distance are structurally similar. Two rows whose signatures are far apart are structurally divergent.

```
lo_fp: bit(48) = '111011100110110011010101010001110111010110100  
11101'
```

THE SCORE VECTOR


Four dimensions. Each is a normalized [0, 1] measurement against the rolling universe. The composite is a fixed weighted combination: 0.40-reconstruction + 0.35-diversity + 0.25-anomaly.

```
lo_score: {  
  composite: 0.859,  
  anomaly: 1.000,  
  reconstruction: 0.894,  
  diversity: 0.815  
}
```

What it replaces

The product, explained — one primitive, fifteen data shapes.

Everything, at once, in 33 seconds.



Eighteen public APIs, processed live, in a single deterministic run.

Real-time. Really real.

Two streaming sources, scored every twenty seconds.

Trade secrets, with timestamps no one can fake.



- The deep technical methods are intended to be protected as trade secrets.
- Every invention is cryptographically time-stamped on a public blockchain.
- Anyone can verify when we created it without ever seeing what it is.